


Participants: Tom Wu, Yuexi Sun, Qinquan Wang, Yitian Kong, Wenhan Dong, Yan Bao, Mike Wang

google docs

link: <https://docs.google.com/document/d/1K5IFrNLaMaNFKR39uyPfVbRyVf2wtXWX1RTmlkS7WU/edit?usp=sharing>

Wednesday meeting question preparation

Stage	What you do	What changes
1. Run a pretrained model	Give it a student paragraph and ask for feedback	Nothing in the model
2. Improve the prompt	Add the rubric, examples, and teaching instructions	The input
3. Supervised fine-tuning	Train on examples of good instructor feedback	Some model parameters
4. Reinforcement learning	Generate feedback, score it, and train toward better responses	Input prompt changed by learning



Emphasize on model and dataset

Tom:

1. What size of model should we go? Where to train?
2. What feedback rubric do we want? What is considered a good training set for fine-tuning? **Rubric** for reinforcement learning (bandit selector)? What rubric guides the selector to choose certain strategies?
3. Evaluator model for implementing the rubric for reinforcement learning? (very specific)

Natalie: Dataset trustworthiness? Ask owner for data? (template for applying dataset)

Real-time feedback or after-submission feedback?

2. Should the SLM help in real time as students write (like Copilot), or should it give feedback after students submit a full draft?

3. What kind of research questions or analysis should the logs support? Are you looking for quantitative data, qualitative data, or both?

Don't overengineering!!!!!!!!!!

How to design system architecture?